

Parallel Processing

Modern Computational Methods in Physics 2.

Balázs *PÁL*

¹Eötvös Loránd University

²Wigner Research Centre for Physics

Questions

- „How can I run this code in parallel?”
- „What should I use in Python to run my code in parallel?”
- „Can I run a Jupyter Notebook in parallel somehow?”
- „Does Python code run in parallel by default?”
- ⋮


Answers

- No. Just no. Please no. No.
- How much is behind „parallelisation”?
- When and why are we even using it?
- In what situations and how can a physicist use it?


Important concepts

- Fundamental terms
 - ▶ Thread
 - ▶ Process
 - ▶ Core
- Computational methods
 - ▶ Serial/Sequential
 - ▶ Parallel
- Parallelisation methods
 - ▶ Multi-threading
 - ▶ Multi-processing
 - ▶ ...

Serial operation schematic diagram



Parallel computing

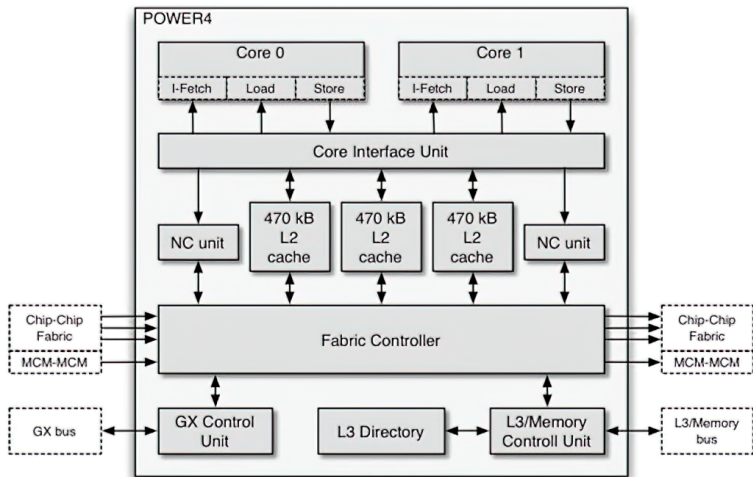


Source: ResearchGate

CPU („Central Processing Unit”) and cores

- „The brain of a computer”, it is responsible for
 - ▶ the calculation of arithmetic operations,
 - ▶ performing logical operations,
 - ▶ operation of the other hardware components.
- Nowadays all CPUs are multi-core
 - ▶ First multi(two)-core model: POWER4 (IBM, 2001)

Threads, Processes, Cores and other creatures



Source: ibm.com

Threads, Processes, Cores and other creatures

CPU („Central Processing Unit”) and cores

- „The brain of a computer”, it is responsible for
 - ▶ the calculation of arithmetic operations,
 - ▶ performing logical operations,
 - ▶ operation of the other hardware components.
- Nowadays all CPUs are multi-core
 - ▶ First multi(two)-core model: POWER4 (IBM, 2001)

Processes, Threads

- Process
 - ▶ Name for a single, running program
- Thread
 - ▶ A running sequence of operations
 - ▶ A process can be broken down into several threads
 - ▶ They always run „in parallel”


Running Threads

Managing threads

- ① On the application level
 - ▶ It is up to the user to decide which part of the program to split into threads and how
- ② On the OS level
 - ▶ The OS thread manager decides the order at which threads are run

Running in parallel

- With **more cores** than threads, they actually run in parallel
- With **more threads** than cores, the CPU periodically cycles through them



Source: Cburnett, Wikipedia

An example for a parallelisable problem

Cosmological N-body simulations


Lots and lots of similar calculations...

$$\Psi(\mathbf{q}, \tau) = \int \frac{d^3 k}{(2\pi)^3} e^{i\mathbf{k} \cdot \mathbf{q}} \frac{i\mathbf{k}}{k^2} \delta_L(\mathbf{k})$$

$$\mathbf{g}_i = -\nabla \varphi(\mathbf{x}_i) = G \sum_{\substack{j=1 \\ j \neq i}}^N m_j \mathcal{F}(\mathbf{x}_i - \mathbf{x}_j)$$

$$\mathbf{x} = \mathbf{q} + \Psi(\mathbf{q}, \tau)$$

$$\dot{\mathbf{x}} = \frac{\dot{D}(a)}{D(1)} \Psi(\mathbf{q}, \tau)$$



But there are many more...



Examples of parallelizable problems

- Monte-Carlo-simulation
- Numerical integration
- Machine learning algorithms (e.g. random forest)


⋮

- Computer visualisations and modelling
- Problems related to image processing
- Discrete Fourier transform


⋮



VS



www.educba.com



Race condition

Parallelisation in practice

Several tools now make programming parallel code much easier:

- Python: threading, multiprocessing etc.
- C++17: Standard libraries
- C/C++/Fortran: OpenMP, CUDA, etc.

Race condition

Parallelisation in practice

Several tools now make programming parallel code much easier:

- Python: threading, multiprocessing etc.
- C++17: Standard libraries
- C/C++/Fortran: OpenMP, CUDA, etc.

Issues that needs to be addressed

- There would be plenty without the tools mentioned above, but there's still many more...
- Probably the most important one is „race condition”
 - ▶ Two threads want to access the same memory space at the same time
 - ▶ A typical case of the dreaded „undefined behaviour”
 - ▶ A critical bug that can completely mess up the behaviour of any running process

Race condition

A simple (but kinda dumb) example

race-condition-serial.cpp

```
int addition(int &x) {
    int new_x = x + 1;
    return new_x;
}

int main(int argc, char const *argv[])
{
    // Starting number
    int x = 0;

    // Adding +1 to it 4 times
    for(int i = 0; i < 4; i++) {
        x = addition(x);
        std::cout << "Current value of `x` is " << x << std::endl;
    }

    return 0;
}
```

race-condition-parallel.cpp

```
int addition(int x)
{
    int new_x = x + 1;
    return new_x;
}

int main(int argc, char const *argv[])
{
    // Starting number
    int x = 0;

    // Parallelization
    std::vector<std::future<int>> future_vec;

    // Adding +1 to it 4 times
    for(int i = 0; i < 4; i++)
    {
        future_vec.push_back(std::async(std::launch::async, addition, x));
    }

    for(int i = 0; i < 4; i++)
    {
        auto new_x = future_vec[i].get();
        std::cout << "Current value of `x` is " << new_x << std::endl;
    }


    return 0;
}
```

```
(march) | (master@sky) ~$ ./race_condition_serial
Current value of `x` is 1
Current value of `x` is 2
Current value of `x` is 3
Current value of `x` is 4
```

```
(march) | (master@sky) ~$ ./race_condition_parallel
Current value of `x` is 1
Current value of `x` is 1
Current value of `x` is 1
Current value of `x` is 1
```

Race condition

Northeast blackout of 2003



Source: clevescene.com

Motivation

How much time can be saved by parallelisation in real cases?

Amdahl's law

Motivation

How much time can be saved by parallelisation in real cases?

- The runtime of a program can be expressed as follows:

$$T = T \cdot S + T \cdot P$$

where $S + P = 1$ and which are usually the ratio of the number of parts of an algorithm that can be run only in a serial manner (S) and the number of parts that can be parallelized (P).

Motivation

How much time can be saved by parallelisation in real cases?

- The runtime of a program can be expressed as follows:

$$T = T \cdot S + T \cdot P$$

where $S + P = 1$ and which are usually the ratio of the number of parts of an algorithm that can be run only in a serial manner (S) and the number of parts that can be parallelized (P).

- If the parallelizable part (P) is divided into several (N) threads, the runtime of the program is reduced:

$$T_{\text{new}} = T \cdot S + T \cdot \frac{P}{N}$$

Motivation

How much time can be saved by parallelisation in real cases?

- The speedup (Q) now can be expressed as:

$$Q(N) = \frac{T}{T_{\text{new}}} = \frac{\cancel{S}}{\cancel{S} \cdot S + \cancel{S} \cdot \frac{P}{N}} = \frac{1}{S + \frac{P}{N}}$$

Amdahl's law

Motivation

How much time can be saved by parallelisation in real cases?

- The speedup (Q) now can be expressed as:


$$Q(N) = \frac{T}{T_{\text{new}}} = \frac{\cancel{T}}{\cancel{T} \cdot S + \cancel{T} \cdot \frac{P}{N}} = \frac{1}{S + \frac{P}{N}}$$

- After some simplification, we get back the usual form of Amdahl's law:

$$Q(S, N) = \frac{1}{S + \frac{P}{N}} = \frac{1}{S + \frac{1-S}{N}},$$

which specifies that the speedup depends solely on the proportion of parts to be run in series and the number of threads.

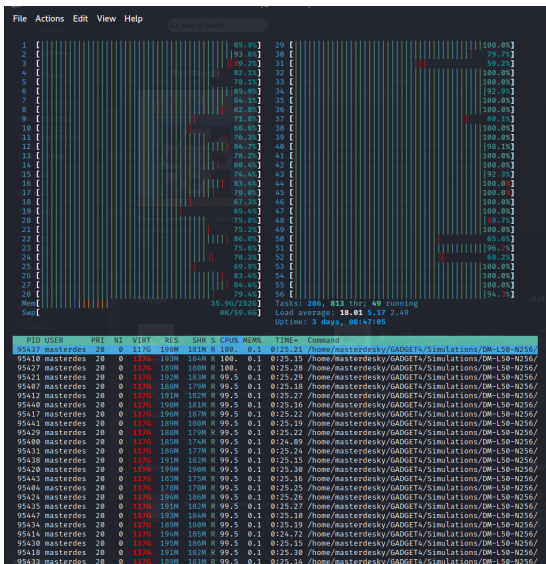
Amdahl's law



Parallelisation in physics

In general

- Running simulation software (e.g. Computer Simulation and Modelling, MSc course – OpenFoam, GADGET4, HOOMD-blue)
- Working on servers (e.g. onco2, tesla, atys, veo1 (known as ‘kooplex’) etc. at ELTE)
- Working on HPC clusters (e.g. ELTE *Atlasz* cluster or the Hungarian Komondor HPC cluster, etc.)



Parallelisation in physics

On servers and clusters

```
#!/bin/bash -l
#SBATCH --job-name="G4M2Norm"
#SBATCH --partition=hpc2019
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=18
#SBATCH --cpus-per-task=1
#SBATCH --time=24:00:00
#SBATCH --exclusive

echo
echo "Running on hosts: $SLURM_NODELIST"
echo "Running on $SLURM_NNODES nodes."
echo "Running on $SLURM_NPROCS processors."
echo

export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/users/lordpb666/opt/gsl-2.6/lib/
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/users/lordpb666/opt/fftw-3.3.9/lib/

mpirun -np $SLURM_NPROCS /users/lordpb666/apps/GADGET4/Gadget4 \
/users/lordpb666/apps/GADGET4/Simulations/DM-N65536-M2-L200-GTF5/GADGET4_GTF5.param
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
314254	hpc2019	7-28-61	hanyecz	R	23-13:29:58	1	cn18-25
314255	hpc2019	7-62-95	hanyecz	R	23-13:29:52	1	cn18-26
314256	hpc2019	7-96-128	hanyecz	R	23-13:29:49	1	cn18-07
314257	hpc2019	14-1-42	hanyecz	R	23-13:26:30	1	cn18-08
314258	hpc2019	14-43-84	hanyecz	R	23-13:25:27	1	cn18-09
314259	hpc2019	14-85-12	hanyecz	R	23-13:24:23	1	cn18-10
316125_8	hpc2019	evo22.sb	skiszkao	R	12-10:50:51	1	cn18-24
325390	hpc2019	atlasz_b	geobarbi	R	5-06:35:54	5	cn18-[04,11-12,16,18]
326521	hpc2019	atlasz_b	geobarbi	R	5-12:35:29	1	cn18-06
326839	hpc2019	atlasz_b	geobarbi	R	5-06:07:18	1	cn18-35
327902	hpc2019	atlasz_b	geobarbi	R	4-04:12:25	2	cn18-[05,27]
327919	hpc2019	atlasz_b	geobarbi	R	3-16:49:20	1	cn18-33
327942	hpc2019	atlasz_b	geobarbi	R	3-10:09:39	1	cn18-03

Parallelisation in physics

Using Python

Prerequisites

- E.g. the threading, multiprocessing, joblib, subprocess etc. packages mentioned in the previous section
- Multiple CPU cores available for ordering (e.g. this is 2 in the case of Kooplex)

```
_ = Parallel(n_jobs=n_jobs)(delayed(create_frame)(i, P,  
N, n_steps, gl,  
outdir, fmt, fdpi) for i, gl in enumerate(tqdm(grid_lims)))
```

```
# Takes lot of time to write all files and ping all sites!  
for target in df_n['URL']:  
    ping_command = 'ping -D -c {0} -i {1} -s {2} {3}'.format(n_packet, interval, packet_size, target)  
    output = "{0}{1}.txt".format(data, target)  
    with open(output, 'w') as f:  
        # Using `Popen` here to run pings "parallel"  
        #print('Pinging {}...'.format(target))  
        process = subprocess.Popen(ping_command.split(' '), stdout=f)
```